

# The Lottery Ticket Hypothesis

Coco Zhang and Noah Andrew

May 8th, 2024



# The Theorem (Frankle & Carbin, 2019)

“Dense, randomly-initialized, feed-forward networks contain subnetworks (winning tickets) that—when trained in isolation—it can match the test accuracy of the original network after training for at most the same number of iterations.”

# Neural network pruning



- **Objective**

- Improve performance & accuracy and decrease storage requirements by reducing the parameter counts of trained networks

- **One-shot Pruning**

1. Randomly initialize a neural network  $f(x; \theta_0)$  (where  $\theta_0 \sim \mathcal{D}_\theta$ ).
2. Train the network for  $j$  iterations, arriving at parameters  $\theta_j$ .
3. Prune  $p\%$  of the parameters in  $\theta_j$ , creating a mask  $m$ .
4. Reset the remaining parameters to their values in  $\theta_0$ , creating the winning ticket  $f(x; m \odot \theta_0)$ .

- **Iterative Pruning**

- Repeats the process of one-shot pruning (prune, reset, retrain) over  $n$  rounds, where
- Each round prunes  $p^{1/n} \%$  of weights that survived the previous round
- More computationally and time intensive, since the network must be pruned and re-trained many times

# Experiments (Frankle & Carbin, 2019)



- Establishing the lottery ticket hypothesis for various neural network architectures.
- Findings:
  - **Iterative pruning**: extracts smaller winning tickets.
  - **One-shot pruning**: identifies winning tickets without repeated training.


Network	Lenet	Conv-2	Conv-4	Conv-6	Resnet-18	VGG-19
				64, 64, pool	16, 3x[16, 16]	2x64 pool 2x128
			64, 64, pool	128, 128, pool	3x[32, 32]	pool, 4x256, pool
Convolutions		64, 64, pool	128, 128, pool	256, 256, pool	3x[64, 64]	4x512, pool, 4x512
FC Layers	300, 100, 10	256, 256, 10	256, 256, 10	256, 256, 10	avg-pool, 10	avg-pool, 10
All/Conv Weights	266K	4.3M / 38K	2.4M / 260K	1.7M / 1.1M	274K / 270K	20.0M
Iterations/Batch	50K / 60	20K / 60	25K / 60	30K / 60	30K / 128	112K / 64
Optimizer	Adam 1.2e-3	Adam 2e-4	Adam 3e-4	Adam 3e-4	← SGD 0.1-0.01-0.001 Momentum 0.9 →	
Pruning Rate	fc20%	conv10% fc20%	conv10% fc20%	conv15% fc20%	conv20% fc0%	conv20% fc0%

# Our Deliverable

- We extended the `training` module from the digitize demo to include the following:

```
def winning_ticket(trained_network: NeuralNetwork, copy_network: NeuralNetwork, amount: float):  
    """  
    Function that takes a trained network and a previously made copy of that network  
    and prunes away the lowest valued parameter weights by l1 norm. The copy network is  
    modified in place to be a "winning ticket," i.e. it retains the random initialization  
    values that later became highest magnitude in the trained network.  
    """
```

# Identifying a “winning ticket” with one-shot pruning



- We applied one-shot pruning to a feedforward neural network on the MNIST dataset.
  - `torch.nn.prune` package
  - `prune.l1_unstructured`: removes a variable percentage of the smallest- magnitude weights in each parameter

# Discussions

- Experiments demonstrate that the function learned by a neural network can often be represented with fewer parameters – a winning ticket exists.
- Winning ticket initialization is key to achieving better test accuracy than the original unpruned neural network, but initialization is also sparsity dependent.
- Limitations:
  - only finding winning tickets through sparse pruning
  - may not be as effective for larger datasets like Imagenet

# Questions or Comments?